

Business Case: Porter: Neural Networks Regression

Defining Problem Statement & Data Import

Problem Statement:

Porter is India's Largest Marketplace for Intra-City Logistics. Leader in the country's \$40 billion intra-city logistics market, Porter strives to improve the lives of 1,50,000+ driver-partners by providing them with consistent earning & independence. Currently, the company has serviced 5+ million customers

Porter works with a wide range of restaurants for delivering their items directly to the people.

Porter has a number of delivery partners available for delivering the food, from various restaurants and wants to get an estimated delivery time that it can provide the customers on the basis of what they are ordering, from where and also the delivery partners.

This dataset has the required data to train a regression model that will do the delivery time estimation, based on all those features.

Dataset:

https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/015/039/original/dataset.csv
1663710760

Data Dictionary:

Each row in this file corresponds to one unique delivery. Each column corresponds to a feature as explained below.

- market_id : integer id for the market where the restaurant lies
- created_at : the timestamp at which the order was placed
- actual_delivery_time : the timestamp when the order was delivered
- store_primary_category : category for the restaurant
- order_protocol : integer code value for order protocol(how the order was placed ie: through porter, call to restaurant, pre booked, third part etc)
- total_items subtotal : final price of the order
- num_distinct_items : the number of distinct items in the order
- min_item_price : price of the cheapest item in the order
- max_item_price : price of the costliest item in order

- total_onshift_partners : number of delivery partners on duty at the time order was placed
- total_busy_partners : number of delivery partners attending to other tasks
- total_outstanding_orders : total number of orders to be fulfilled at the moment

Analysing basic metrics

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: #Loading of dataset
df = pd.read_csv("../Hemangi/dataset.csv")
df.head()
```

```
Out[2]:
```

| | market_id | created_at | actual_delivery_time | store_id | store_id |
|---|-----------|---------------------|----------------------|----------------------------------|----------|
| 0 | 1.0 | 2015-02-06 22:24:17 | 2015-02-06 23:27:16 | df263d996281d984952c07998dc54358 | |
| 1 | 2.0 | 2015-02-10 21:49:25 | 2015-02-10 22:56:29 | f0ade77b43923b38237db569b016ba25 | |
| 2 | 3.0 | 2015-01-22 20:39:28 | 2015-01-22 21:09:09 | f0ade77b43923b38237db569b016ba25 | |
| 3 | 3.0 | 2015-02-03 21:21:45 | 2015-02-03 22:13:00 | f0ade77b43923b38237db569b016ba25 | |
| 4 | 3.0 | 2015-02-15 02:40:36 | 2015-02-15 03:20:26 | f0ade77b43923b38237db569b016ba25 | |

```
In [3]: df.shape #to observe shape of data
```

```
Out[3]: (197428, 14)
```

- Dataset is of 205843 rows and 7 attributes.

```
In [4]: df.info() #to observe the data type
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null  float64
1   created_at                           197428 non-null  object
2   actual_delivery_time                 197421 non-null  object
3   store_id                             197428 non-null  object
4   store_primary_category               192668 non-null  object
5   order_protocol                       196433 non-null  float64
6   total_items                          197428 non-null  int64
7   subtotal                             197428 non-null  int64
8   num_distinct_items                  197428 non-null  int64
9   min_item_price                      197428 non-null  int64
10  max_item_price                      197428 non-null  int64
11  total_onshift_partners               181166 non-null  float64
12  total_busy_partners                  181166 non-null  float64
13  total_outstanding_orders             181166 non-null  float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB

```

Check for Duplicate Values

```
In [5]: df.duplicated().sum()
```

```
Out[5]: 0
```

- There are no duplicate instances in the data

```
In [6]: df.dropna(inplace=True)
df.shape
```

```
Out[6]: (176248, 14)
```

Check for Missing Values

```
In [7]: # Check Missing Values
df.isna().sum()
```

```

Out[7]: market_id                0
        created_at              0
        actual_delivery_time     0
        store_id                0
        store_primary_category   0
        order_protocol           0
        total_items              0
        subtotal                 0
        num_distinct_items       0
        min_item_price           0
        max_item_price           0
        total_onshift_partners   0
        total_busy_partners      0
        total_outstanding_orders 0
        dtype: int64

```

Preprocessing & Feature Creation

```
In [8]: #Setting datetime format for respective columns
df['created_at'] = pd.to_datetime(df['created_at'], dayfirst=True)
df['actual_delivery_time'] = pd.to_datetime(df['actual_delivery_time'], dayfirst=
```

```
In [9]: #Creating new features from Date Time
df['order_hour'] = df['created_at'].dt.hour
df['order_day'] = df['created_at'].dt.dayofweek
```

```
In [10]: #Creating Target Variable & converting it to number of minutes
df['time_taken'] = df['actual_delivery_time'] - df['created_at']
df['delivery_time_taken'] = pd.to_timedelta(df['time_taken']) / pd.Timedelta('60s')
```

```
In [11]: #Dropping the columns that are no longer required
df.drop(['time_taken', 'created_at', 'actual_delivery_time', 'store_id'], axis=1, inp
```

```
In [12]: df['store_primary_category'] = df['store_primary_category'].astype('category').cat
```

```
In [13]: #Converting required float columns to int datatype
float_cols = ['total_onshift_partners', 'total_busy_partners', 'total_outstanding_
df[float_cols] = df[float_cols].astype('int64')
```

```
In [14]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 176248 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             176248 non-null  int64
1   store_primary_category                 176248 non-null  int8
2   order_protocol                         176248 non-null  int64
3   total_items                           176248 non-null  int64
4   subtotal                              176248 non-null  int64
5   num_distinct_items                    176248 non-null  int64
6   min_item_price                        176248 non-null  int64
7   max_item_price                        176248 non-null  int64
8   total_onshift_partners                 176248 non-null  int64
9   total_busy_partners                   176248 non-null  int64
10  total_outstanding_orders               176248 non-null  int64
11  order_hour                             176248 non-null  int64
12  order_day                              176248 non-null  int64
13  delivery_time_taken                   176248 non-null  int64
dtypes: int64(13), int8(1)
memory usage: 19.0 MB
```

```
In [15]: df.describe().T
```

Out[15]:

| | count | mean | std | min | 25% | 50% | 75% | 90% |
|---------------------------------|----------|-------------|-------------|-------|--------|--------|--------|--------|
| market_id | 176248.0 | 2.743747 | 1.330911 | 1.0 | 2.0 | 2.0 | 4.0 | 5.0 |
| store_primary_category | 176248.0 | 35.891482 | 20.728572 | 0.0 | 18.0 | 38.0 | 55.0 | 61.0 |
| order_protocol | 176248.0 | 2.911687 | 1.512920 | 1.0 | 1.0 | 3.0 | 4.0 | 5.0 |
| total_items | 176248.0 | 3.204592 | 2.673899 | 1.0 | 2.0 | 3.0 | 4.0 | 6.0 |
| subtotal | 176248.0 | 2696.498939 | 1828.922584 | 0.0 | 1408.0 | 2221.0 | 3407.0 | 4970.0 |
| num_distinct_items | 176248.0 | 2.674589 | 1.625558 | 1.0 | 1.0 | 2.0 | 3.0 | 5.0 |
| min_item_price | 176248.0 | 684.937730 | 519.911425 | -86.0 | 299.0 | 595.0 | 942.0 | 1295.0 |
| max_item_price | 176248.0 | 1159.886994 | 560.784510 | 0.0 | 799.0 | 1095.0 | 1395.0 | 1795.0 |
| total_onshift_partners | 176248.0 | 44.905276 | 34.529394 | -4.0 | 17.0 | 37.0 | 65.0 | 98.0 |
| total_busy_partners | 176248.0 | 41.845434 | 32.154573 | -5.0 | 15.0 | 35.0 | 62.0 | 90.0 |
| total_outstanding_orders | 176248.0 | 58.206800 | 52.708344 | -6.0 | 17.0 | 41.0 | 85.0 | 140.0 |
| order_hour | 176248.0 | 8.493872 | 8.681474 | 0.0 | 2.0 | 3.0 | 19.0 | 21.0 |
| order_day | 176248.0 | 3.221563 | 2.041332 | 0.0 | 1.0 | 3.0 | 5.0 | 6.0 |
| delivery_time_taken | 176248.0 | 47.271992 | 27.656174 | 1.0 | 35.0 | 44.0 | 56.0 | 70.0 |

In [16]: `df.describe(include = np.number, percentiles=[.25,.5,.75,.90,.95, .99, .999]).round(2)`

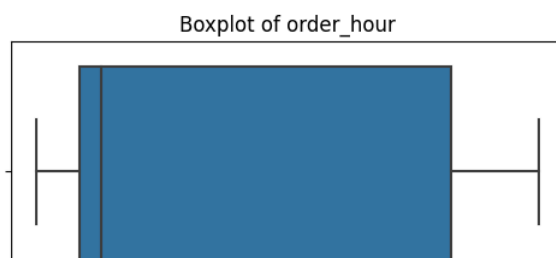
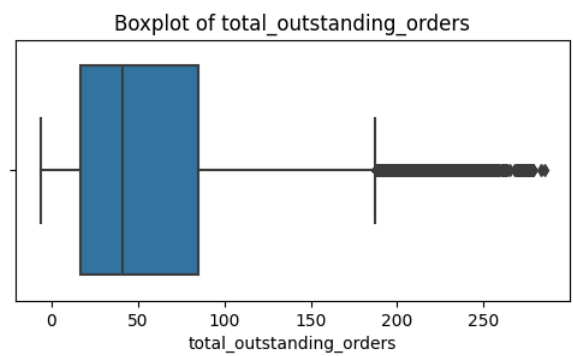
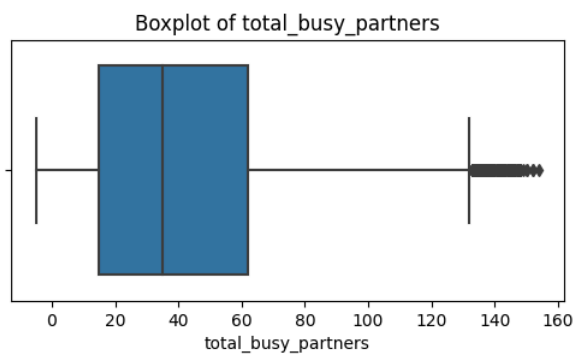
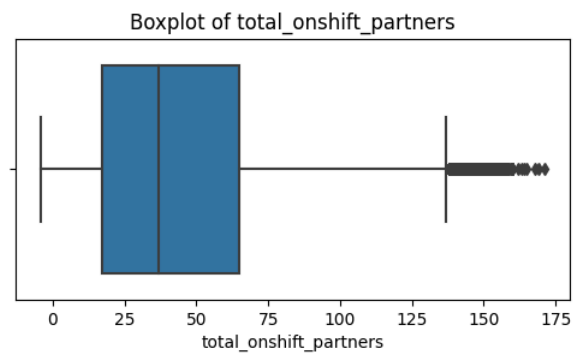
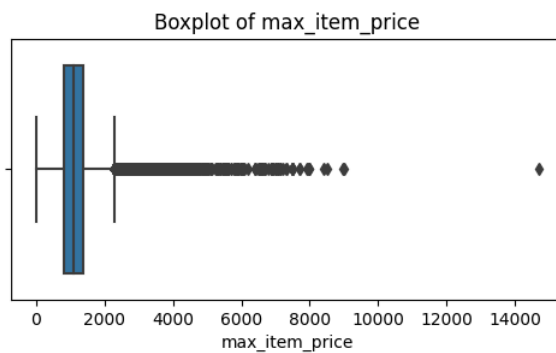
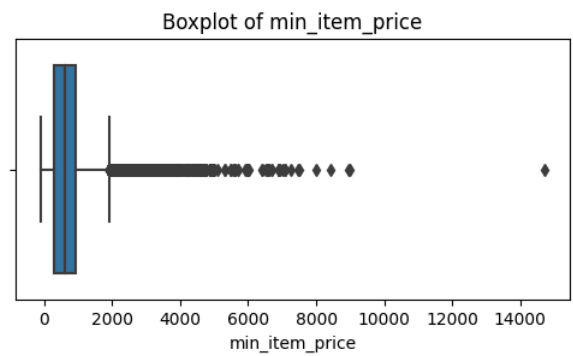
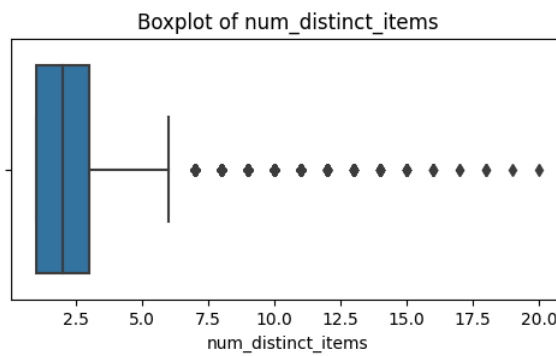
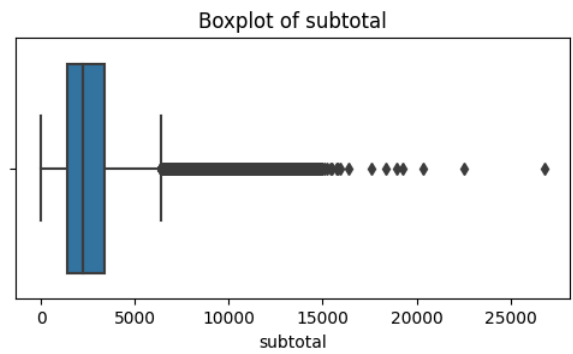
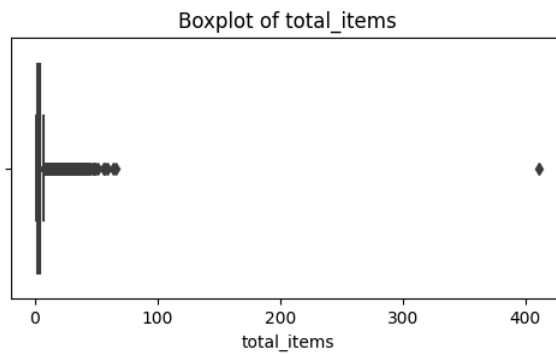
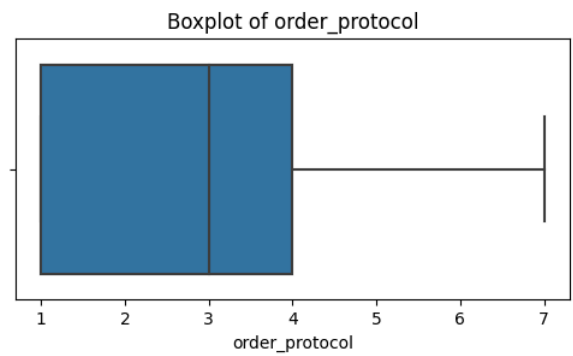
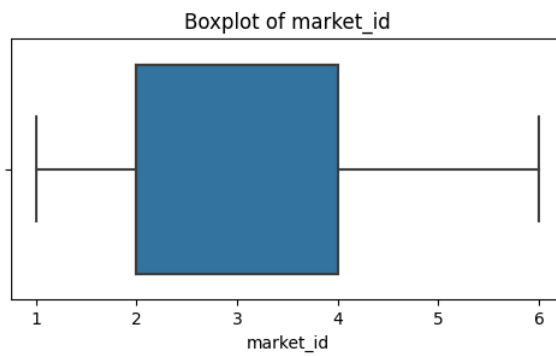
Out[16]:

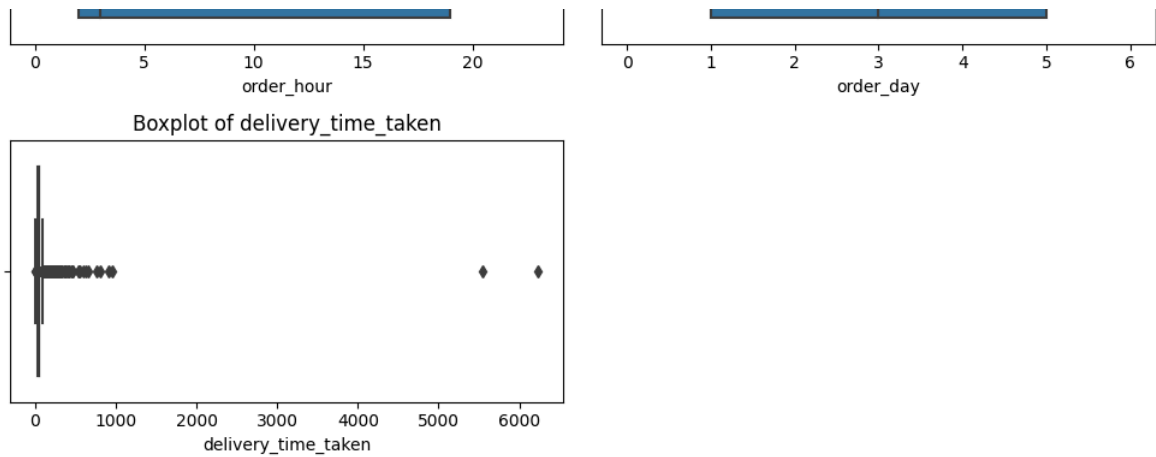
| | count | mean | std | min | 25% | 50% | 75% | 90% |
|---------------------------------|----------|---------|---------|-------|--------|--------|--------|--------|
| market_id | 176248.0 | 2.74 | 1.33 | 1.0 | 2.0 | 2.0 | 4.0 | 5.0 |
| store_primary_category | 176248.0 | 35.89 | 20.73 | 0.0 | 18.0 | 38.0 | 55.0 | 61.0 |
| order_protocol | 176248.0 | 2.91 | 1.51 | 1.0 | 1.0 | 3.0 | 4.0 | 5.0 |
| total_items | 176248.0 | 3.20 | 2.67 | 1.0 | 2.0 | 3.0 | 4.0 | 6.0 |
| subtotal | 176248.0 | 2696.50 | 1828.92 | 0.0 | 1408.0 | 2221.0 | 3407.0 | 4970.0 |
| num_distinct_items | 176248.0 | 2.67 | 1.63 | 1.0 | 1.0 | 2.0 | 3.0 | 5.0 |
| min_item_price | 176248.0 | 684.94 | 519.91 | -86.0 | 299.0 | 595.0 | 942.0 | 1295.0 |
| max_item_price | 176248.0 | 1159.89 | 560.78 | 0.0 | 799.0 | 1095.0 | 1395.0 | 1795.0 |
| total_onshift_partners | 176248.0 | 44.91 | 34.53 | -4.0 | 17.0 | 37.0 | 65.0 | 98.0 |
| total_busy_partners | 176248.0 | 41.85 | 32.15 | -5.0 | 15.0 | 35.0 | 62.0 | 90.0 |
| total_outstanding_orders | 176248.0 | 58.21 | 52.71 | -6.0 | 17.0 | 41.0 | 85.0 | 140.0 |
| order_hour | 176248.0 | 8.49 | 8.68 | 0.0 | 2.0 | 3.0 | 19.0 | 21.0 |
| order_day | 176248.0 | 3.22 | 2.04 | 0.0 | 1.0 | 3.0 | 5.0 | 6.0 |
| delivery_time_taken | 176248.0 | 47.27 | 27.66 | 1.0 | 35.0 | 44.0 | 56.0 | 70.0 |

Outlier Detection

```
In [17]: #Distribution of numerical variables
num_vars = df.select_dtypes('int').columns.tolist()
fig = plt.figure(figsize=(10,21))
i=1
for col in num_vars:
    ax = plt.subplot(7,2,i)
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    i += 1

plt.tight_layout()
plt.show()
```





- Here we can see that many columns have outliers. Lets remove the outliers using IQR method.

```
In [18]: # Remove outliers through IQR method
num_vars = df.select_dtypes('int').columns.tolist()

for col in num_vars:
    # Calculate Q1 and Q3
    q1 = np.percentile(df[col], 25)
    q3 = np.percentile(df[col], 75)

    # Calculate IQR
    iqr = q3 - q1

    # Define lower and upper bounds
    lower_bound = q1 - (1.5 * iqr)
    upper_bound = q3 + (1.5 * iqr)

    # Remove outliers
    df = df.loc[~((df[col] < lower_bound) | (df[col] > upper_bound))]

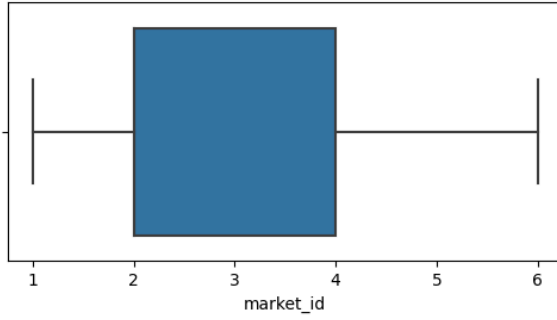
df.shape
```

Out[18]: (142029, 14)

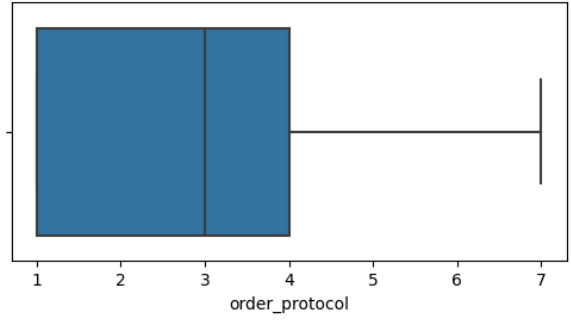
```
In [19]: #Distribution of numerical variables
num_vars = df.select_dtypes('int').columns.tolist()
fig = plt.figure(figsize=(10,21))
i=1
for col in num_vars:
    ax = plt.subplot(7,2,i)
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    i += 1

plt.tight_layout()
plt.show()
```

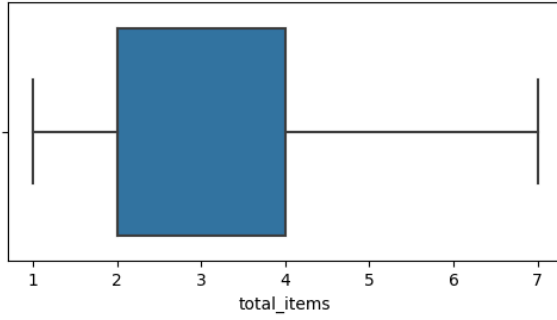

Boxplot of market_id



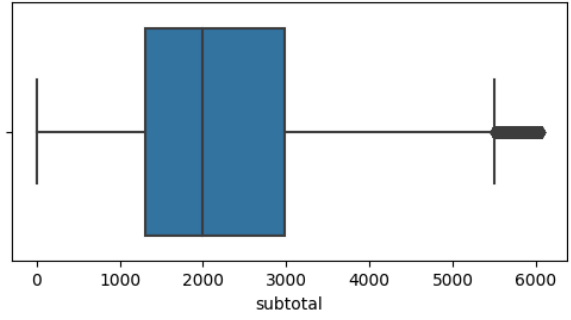
Boxplot of order_protocol



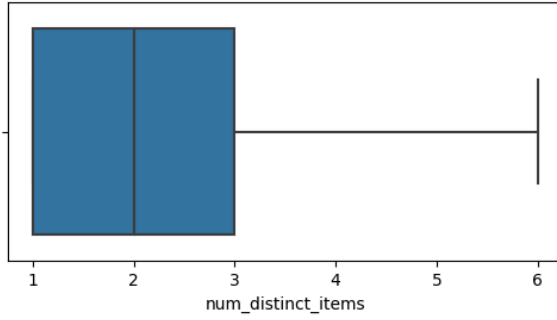
Boxplot of total_items



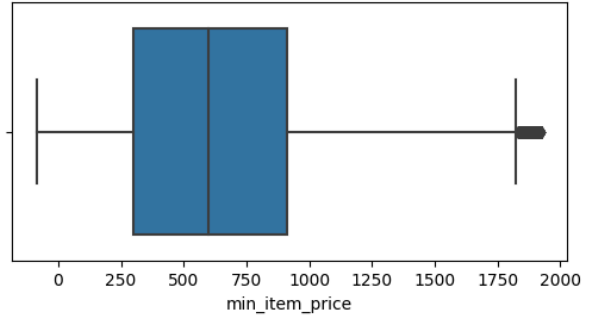
Boxplot of subtotal



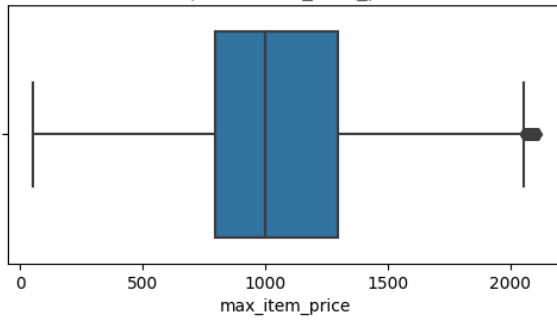
Boxplot of num_distinct_items



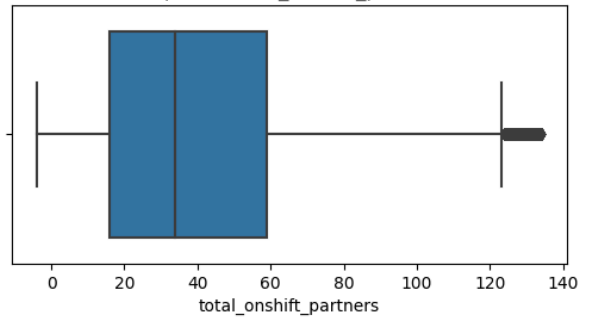
Boxplot of min_item_price



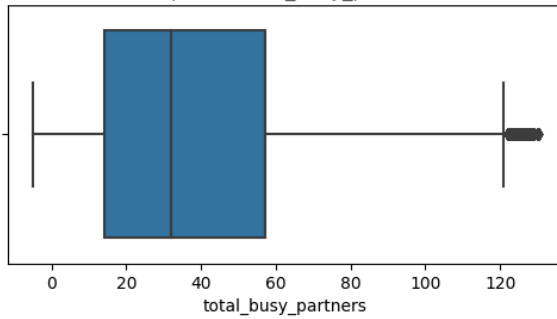
Boxplot of max_item_price



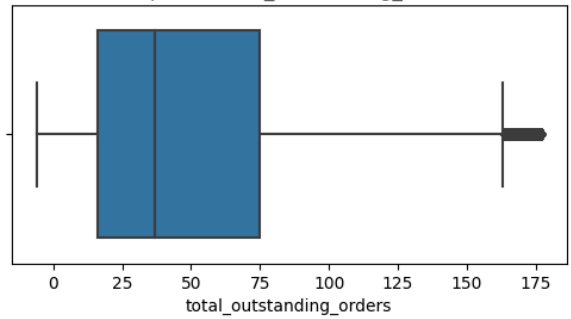
Boxplot of total_onshift_partners



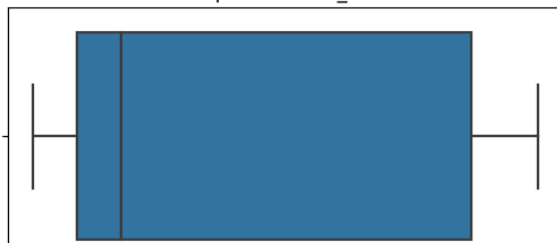
Boxplot of total_busy_partners



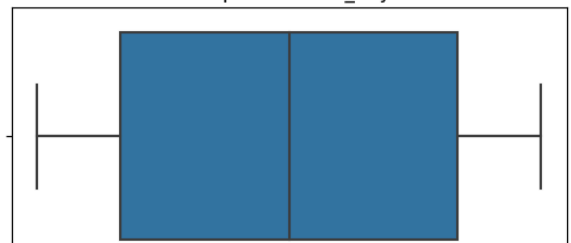
Boxplot of total_outstanding_orders

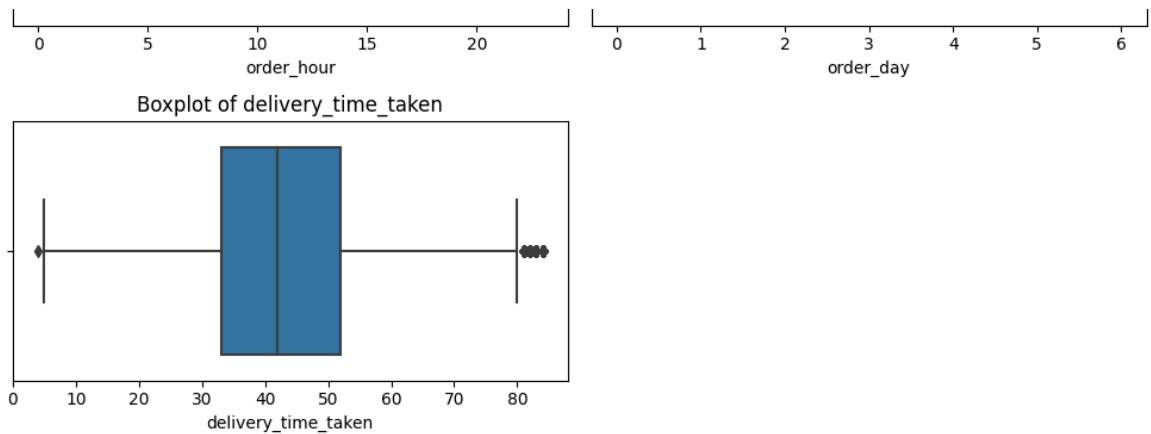


Boxplot of order_hour



Boxplot of order_day





Non-Graphical Analysis

```
In [20]: # Replace min & max item price with average price
df['avg_item_price'] = (df['min_item_price'] + df['max_item_price']) / 2
df.drop(columns=['min_item_price', 'max_item_price'], inplace=True)
```

```
In [21]: # Number of unique values in all columns
unique_num = df.columns.tolist()
for col in unique_num:
    print(f"No. of unique values in {col}: {df[col].nunique()}")
```

```
No. of unique values in market_id: 6
No. of unique values in store_primary_category: 72
No. of unique values in order_protocol: 7
No. of unique values in total_items: 7
No. of unique values in subtotal: 5195
No. of unique values in num_distinct_items: 6
No. of unique values in total_onshift_partners: 139
No. of unique values in total_busy_partners: 136
No. of unique values in total_outstanding_orders: 184
No. of unique values in order_hour: 18
No. of unique values in order_day: 7
No. of unique values in delivery_time_taken: 81
No. of unique values in avg_item_price: 3057
```

```
In [22]: # unique value market_id column (listed in %)
market_id = df['market_id'].value_counts(normalize=True).map(lambda calc: round(
market_id.columns = ['market_id', 'Count']
market_id
```

```
Out[22]:
```

| | market_id | Count |
|---|-----------|-------|
| 0 | 2 | 29.60 |
| 1 | 4 | 25.55 |
| 2 | 1 | 21.22 |
| 3 | 3 | 12.67 |
| 4 | 5 | 10.59 |
| 5 | 6 | 0.37 |

```
In [23]: # unique value order_protocol column(Listed in %)
order_protocol = df['order_protocol'].value_counts(normalize=True).map(lambda ca
order_protocol.columns = ['order_protocol', 'Count']
order_protocol
```

```
Out[23]:
```

| | order_protocol | Count |
|---|----------------|-------|
| 0 | 1 | 27.68 |
| 1 | 3 | 26.46 |
| 2 | 5 | 23.89 |
| 3 | 2 | 11.98 |
| 4 | 4 | 9.60 |
| 5 | 6 | 0.38 |
| 6 | 7 | 0.01 |

```
In [24]: # unique value total_items column(Listed in %)
total_items = df['total_items'].value_counts(normalize=True).map(lambda calc: rc
total_items.columns = ['total_items', 'Count']
total_items
```

```
Out[24]:
```

| | total_items | Count |
|---|-------------|-------|
| 0 | 2 | 30.60 |
| 1 | 3 | 21.74 |
| 2 | 1 | 21.40 |
| 3 | 4 | 13.63 |
| 4 | 5 | 7.02 |
| 5 | 6 | 3.96 |
| 6 | 7 | 1.65 |

```
In [25]: # unique value num_distinct_items column(Listed in %)
num_distinct_items = df['num_distinct_items'].value_counts(normalize=True).map(1
num_distinct_items.columns = ['num_distinct_items', 'Count']
num_distinct_items
```

```
Out[25]:
```

| | num_distinct_items | Count |
|---|--------------------|-------|
| 0 | 2 | 32.72 |
| 1 | 1 | 26.65 |
| 2 | 3 | 22.56 |
| 3 | 4 | 11.58 |
| 4 | 5 | 4.88 |
| 5 | 6 | 1.61 |

```
In [26]: # unique value order_day column(listed in %)
order_day = df['order_day'].value_counts(normalize=True).map(lambda calc: round(
order_day.columns = ['order_day', 'Count']
order_day
```

```
Out[26]:
```

| | order_day | Count |
|---|-----------|-------|
| 0 | 6 | 16.52 |
| 1 | 5 | 15.65 |
| 2 | 4 | 14.64 |
| 3 | 3 | 13.60 |
| 4 | 0 | 13.31 |
| 5 | 2 | 13.27 |
| 6 | 1 | 13.01 |

```
In [27]: df.describe().loc[['min', 'max']]
```

```
Out[27]:
```

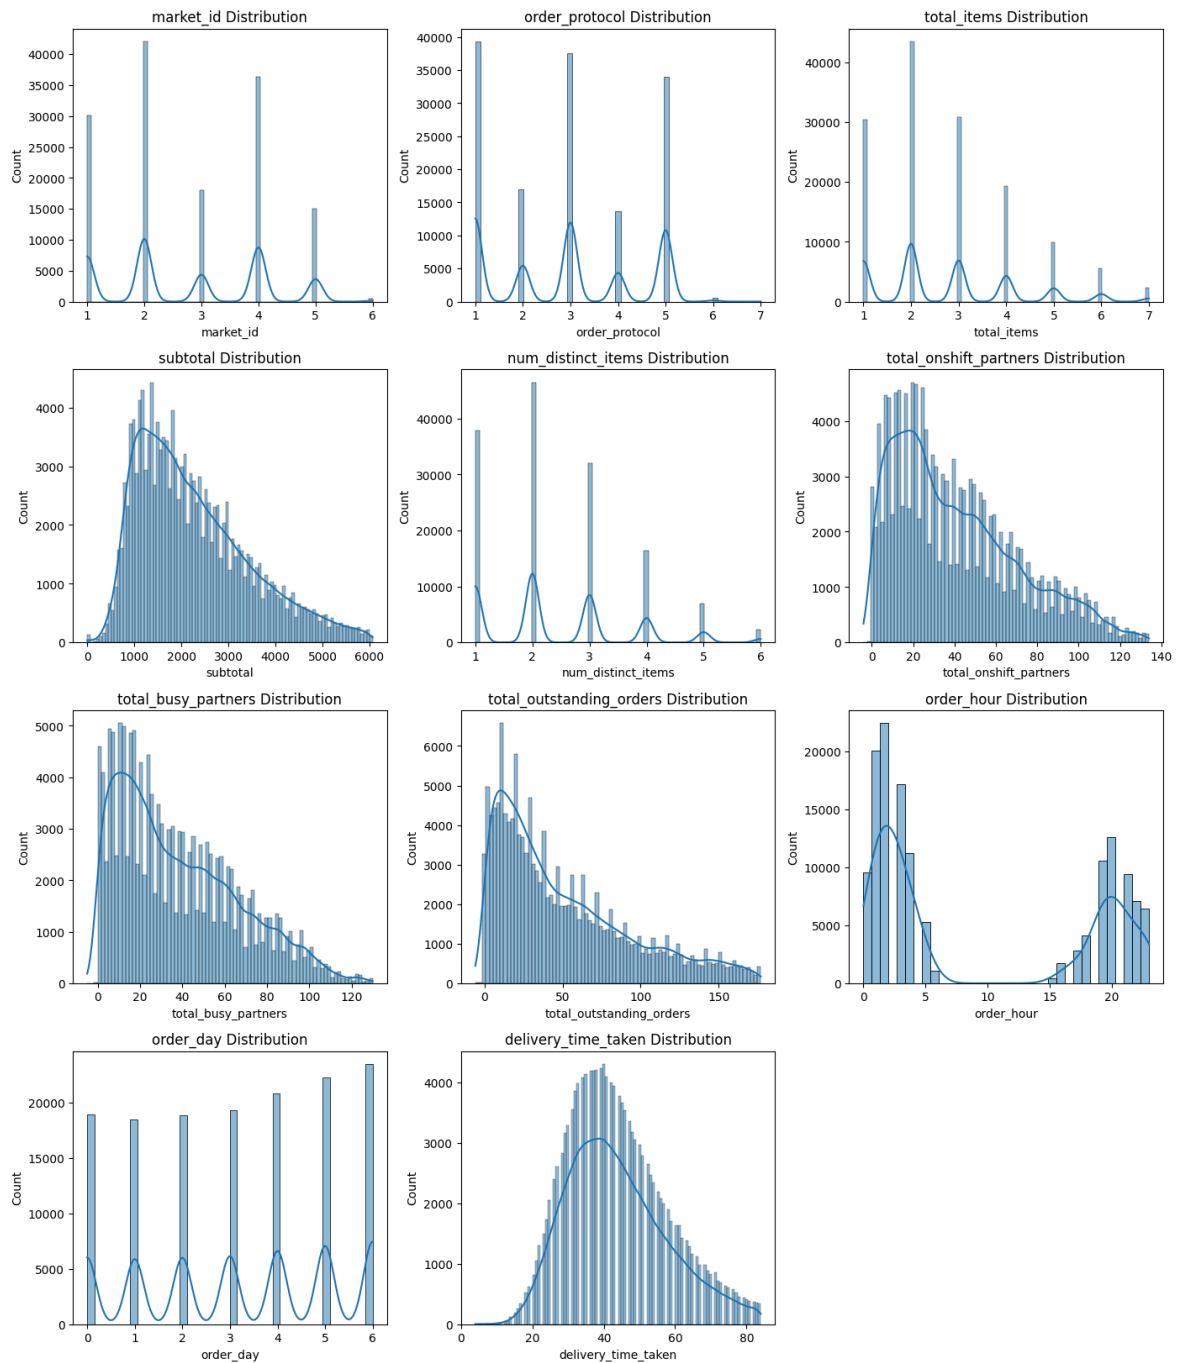
| | market_id | store_primary_category | order_protocol | total_items | subtotal | num_dist |
|-----|-----------|------------------------|----------------|-------------|----------|----------|
| min | 1.0 | 0.0 | 1.0 | 1.0 | 0.0 | |
| max | 6.0 | 72.0 | 7.0 | 7.0 | 6065.0 | |

Univariate Analysis

```
In [28]: plot = df.select_dtypes('int').columns.tolist()

plt.figure(figsize=(14, 20))
i = 1
for col in plot:
    ax = plt.subplot(5, 3, i)
    sns.histplot(data=df, x=col, kde=True)
    plt.title(f'{col} Distribution')
    i += 1

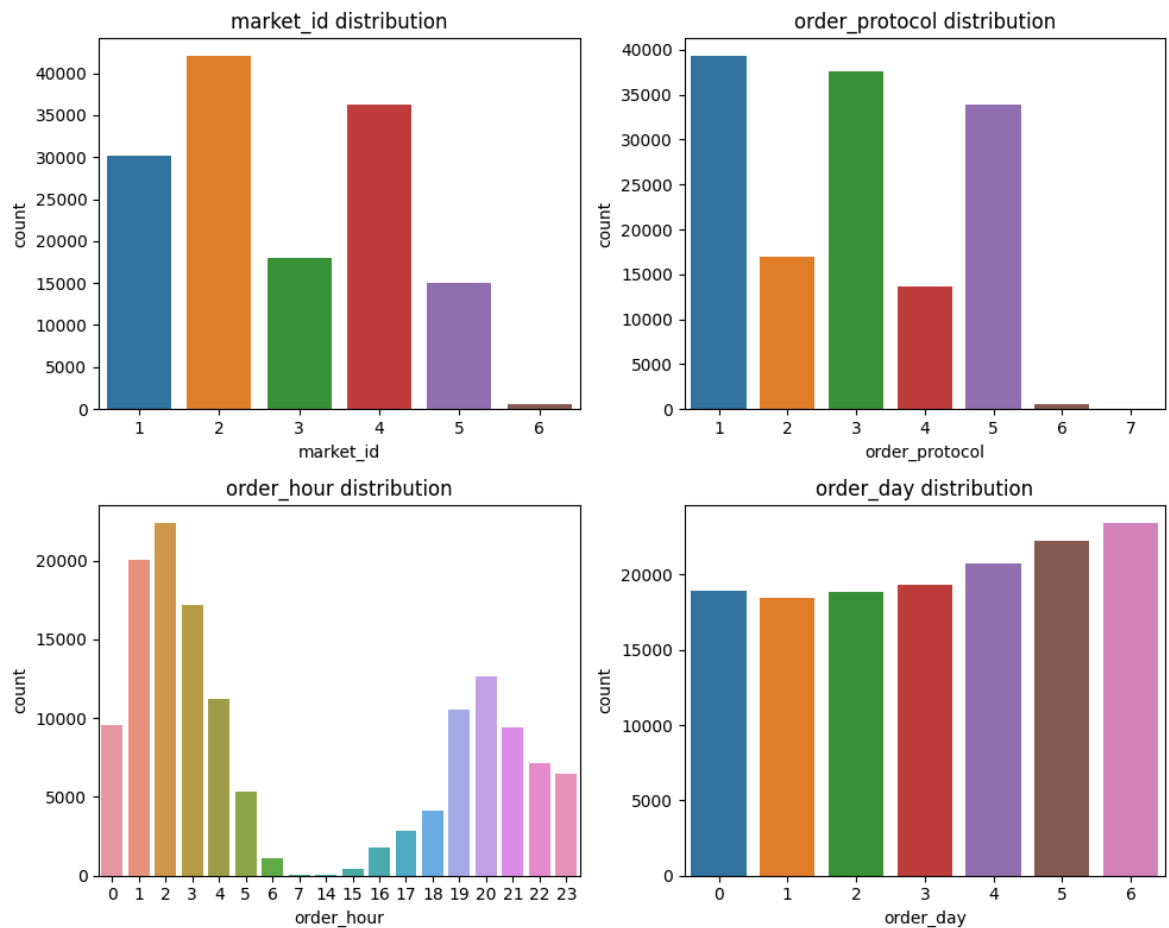
plt.tight_layout()
plt.show()
```



```
In [29]: plot = ['market_id', 'order_protocol', 'order_hour', 'order_day']
```

```
plt.figure(figsize=(10,8))
i=1
for col in plot:
    ax=plt.subplot(2,2,i)
    sns.countplot(x=df[col])
    plt.title(f'{col} distribution')
    i += 1

plt.tight_layout()
plt.show()
```



In [30]: *# Boxplots to analyse the relationship between categorical variables and delivery*

```
cat_cols = ['market_id', 'order_protocol', 'order_hour', 'order_day']
```

```
plt.figure(figsize=(10,16))
```

```
i=1
```

```
for col in cat_cols:
```

```
    ax = plt.subplot(4,1,i)
```

```
    sns.boxplot(data = df, x=col, y='delivery_time_taken')
```

```
    plt.title(f"Impact of {col} on Delivery Time")
```

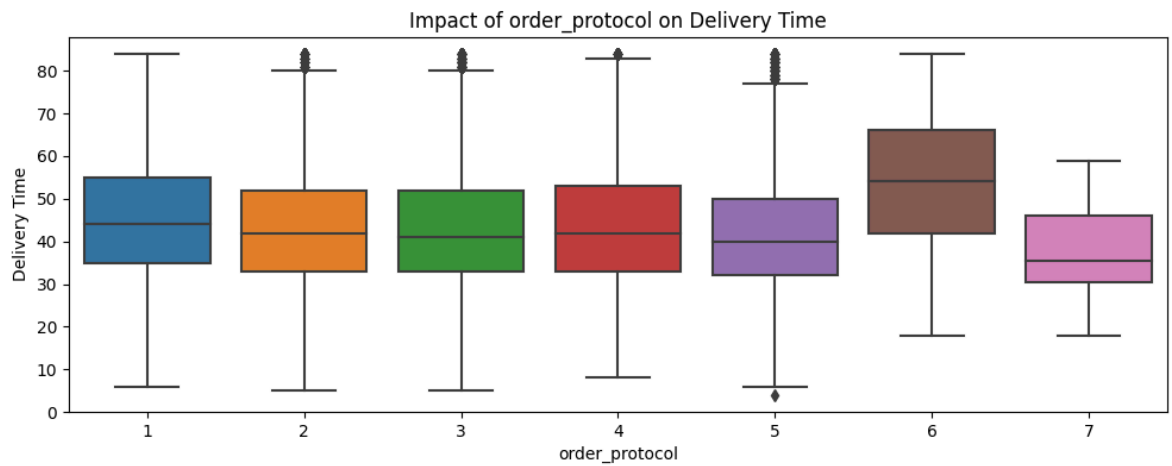
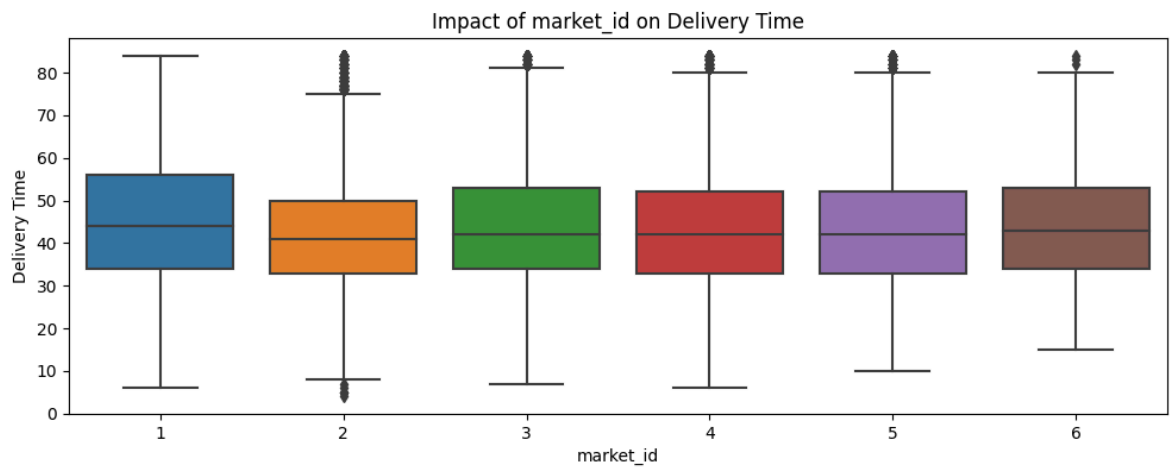
```
    plt.xlabel(col)
```

```
    plt.ylabel('Delivery Time')
```

```
    i+=1
```

```
plt.tight_layout()
```

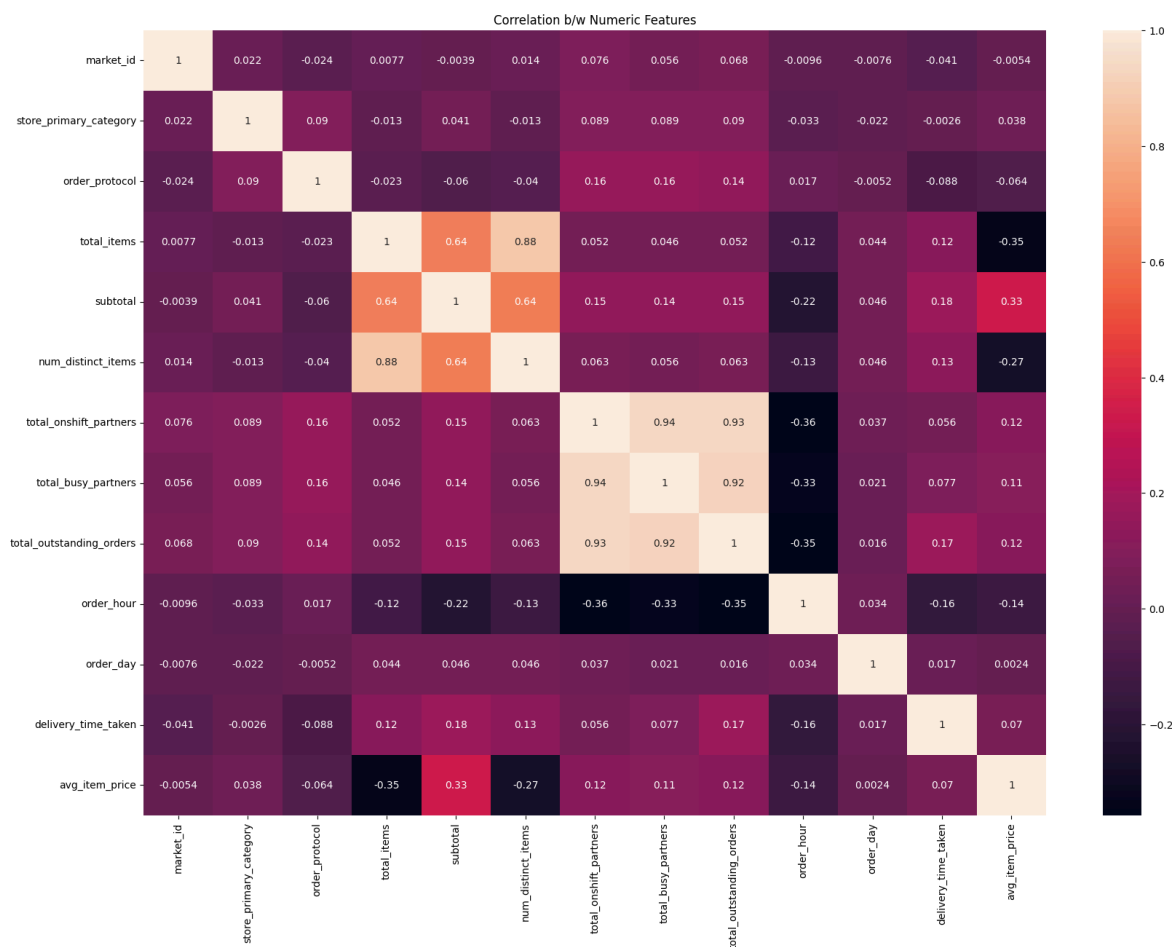
```
plt.show()
```



- Market ID shows no notable impact on delivery duration.

- Orders processed through protocol number 6 show relatively extended delivery times.
- Delivery times are significantly higher at 2 pm than at any other hour.
- Delivery times fluctuate more on Mondays, Saturdays, and Sundays, potentially due to increased order numbers.

```
In [31]: ##Correlation Matrix
corr = df.corr()
plt.figure(figsize=(20,14))
sns.heatmap(corr, annot = True)
plt.title('Correlation b/w Numeric Features')
plt.show()
```



Model Building

```
In [32]: from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import (r2_score, mean_squared_error, mean_absolute_error, mean_squared_error)
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestRegressor
```

```
In [33]: X = df.drop(columns=['delivery_time_taken'])
y = df['delivery_time_taken']
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

```
In [34]: print(f'Shape of x_train: {X_train.shape}')  
print(f'Shape of x_test: {X_test.shape}')
```

Shape of x_train: (113623, 12)
Shape of x_test: (28406, 12)

```
In [35]: print(f'Shape of y_train: {y_train.shape}')  
print(f'Shape of y_test: {y_test.shape}')
```

Shape of y_train: (113623,)
Shape of y_test: (28406,)

Linear Regression

```
In [36]: #Initialising object of Class LinearRegression()  
lr_Test = LinearRegression()    # training LinearRegression model  
lr_Test.fit(X_train,y_train)
```

```
Out[36]: ▾ LinearRegression  
LinearRegression()
```

```
In [37]: #r2 score on train data  
r2_score(y_train,lr_Test.predict(X_train))
```

Out[37]: 0.1674240826950587

```
In [38]: #r2 score on test data  
r2_score(y_test,lr_Test.predict(X_test))
```

Out[38]: 0.16141248202773462

```
In [39]: def MAPE(Y_actual,Y_Predicted):  
    mape = np.mean(np.abs((Y_actual - Y_Predicted)/Y_actual))*100  
    return mape
```

```
In [40]: y_pred = lr_Test.predict(X_test)  
print("MSE:",mean_squared_error(y_test,y_pred)) #MSE  
print("RMSE:",np.sqrt(mean_squared_error(y_test,y_pred))) #RMSE  
print("MAE :",mean_absolute_error(y_test,y_pred) ) #MAE  
print("MAPE :",MAPE(y_test,y_pred) ) #MAPE  
print("r2_score:",r2_score(y_test,y_pred)) #r2score
```

MSE: 166.17941324668996
RMSE: 12.89105943073299
MAE : 10.254421327002385
MAPE : 26.342697198065252
r2_score: 0.16141248202773462

Random Forest Regressor

```
In [41]: regressor=RandomForestRegressor()  
regressor.fit(X_train,y_train)
```

Out[41]: ▾ RandomForestRegressor

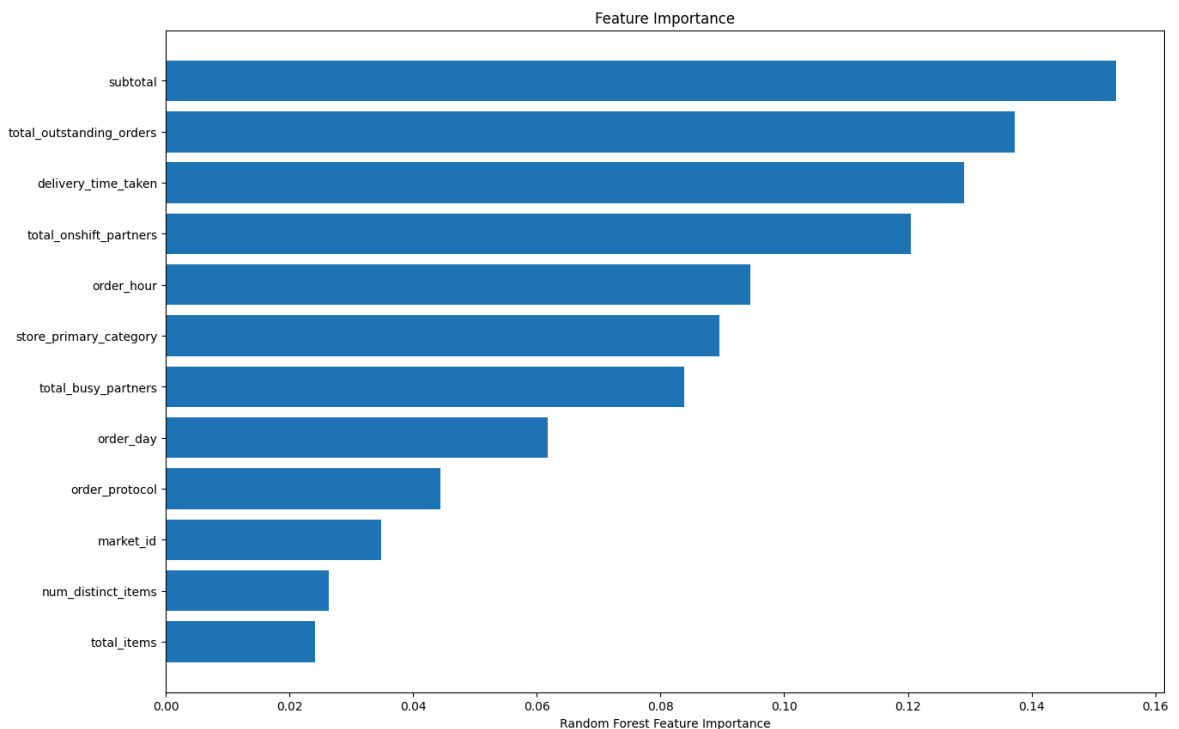
RandomForestRegressor()

```
In [42]: y_pred = regressor.predict(X_test)
print("MSE:", mean_squared_error(y_test, y_pred)) #MSE
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred))) #RMSE
print("MAE :", mean_absolute_error(y_test, y_pred) ) #MAE
print("MAPE :", MAPE(y_test, y_pred) ) #MAPE
print("r2_score:", r2_score(y_test, y_pred)) #r2score
```

MSE: 154.1705630354631
RMSE: 12.416543924758736
MAE : 9.849843166936564
MAPE : 25.2465187704979
r2_score: 0.22201247871555574


Feature Importance


```
In [43]: sorted_idx = regressor.feature_importances_.argsort()
plt.figure(figsize=(15,10))
plt.barh(df.columns[sorted_idx], regressor.feature_importances_[sorted_idx])
plt.title('Feature Importance')
plt.xlabel("Random Forest Feature Importance")
plt.show()
```





```
In [44]: from sklearn import preprocessing
scaler = preprocessing.MinMaxScaler()
x_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2,
```


```
In [45]: import tensorflow as tf
from tensorflow.keras import Model
from tensorflow.keras import Sequential
```



Epoch 1/50
1421/1421  **25s** 10ms/step - loss: 309.3676 - mae: 13.2066 - val_loss: 164.8390 - val_mae: 10.1866


Epoch 2/50
1421/1421  **13s** 9ms/step - loss: 170.8811 - mae: 10.3742 - val_loss: 163.9697 - val_mae: 10.2849


Epoch 3/50
1421/1421  **12s** 9ms/step - loss: 169.9551 - mae: 10.3621 - val_loss: 161.5222 - val_mae: 10.1004


Epoch 4/50
1421/1421  **13s** 9ms/step - loss: 169.4972 - mae: 10.3421 - val_loss: 160.8855 - val_mae: 10.0460


Epoch 5/50
1421/1421  **13s** 9ms/step - loss: 166.3934 - mae: 10.2174 - val_loss: 165.4484 - val_mae: 9.9875


Epoch 6/50
1421/1421  **14s** 10ms/step - loss: 167.0592 - mae: 10.2488 - val_loss: 160.1856 - val_mae: 10.1329


Epoch 7/50
1421/1421  **14s** 10ms/step - loss: 165.5129 - mae: 10.2045 - val_loss: 167.8309 - val_mae: 10.5551


Epoch 8/50
1421/1421  **14s** 10ms/step - loss: 165.2750 - mae: 10.1992 - val_loss: 158.2682 - val_mae: 10.0260


Epoch 9/50
1421/1421  **13s** 9ms/step - loss: 163.7004 - mae: 10.1437 - val_loss: 165.3157 - val_mae: 9.9958


Epoch 10/50
1421/1421  **14s** 10ms/step - loss: 165.2803 - mae: 10.2051 - val_loss: 158.6365 - val_mae: 9.9117


Epoch 11/50
1421/1421  **15s** 11ms/step - loss: 163.3475 - mae: 10.1465 - val_loss: 156.7323 - val_mae: 9.8632


Epoch 12/50
1421/1421  **13s** 9ms/step - loss: 162.2817 - mae: 10.1012 - val_loss: 155.6695 - val_mae: 9.9116


Epoch 13/50
1421/1421  **12s** 9ms/step - loss: 162.1728 - mae: 10.0948 - val_loss: 156.0915 - val_mae: 9.8097


Epoch 14/50
1421/1421  **15s** 10ms/step - loss: 160.3668 - mae: 10.0241 - val_loss: 158.0540 - val_mae: 10.0579


Epoch 15/50
1421/1421  **14s** 10ms/step - loss: 161.5063 - mae: 10.0671 - val_loss: 157.4915 - val_mae: 9.8663

Epoch 16/50
1421/1421  **12s** 8ms/step - loss: 160.1142 - mae: 10.0001 - val_loss: 157.5523 - val_mae: 9.8296

Epoch 17/50
1421/1421  **13s** 9ms/step - loss: 162.1541 - mae: 10.0766 - val_loss: 155.4148 - val_mae: 9.9064

Epoch 18/50
1421/1421  **14s** 10ms/step - loss: 160.3498 - mae: 10.0404 - val_loss: 156.5645 - val_mae: 9.8258

Epoch 19/50
1421/1421  **14s** 9ms/step - loss: 161.1077 - mae: 10.0246 - val_loss: 154.6116 - val_mae: 9.8697

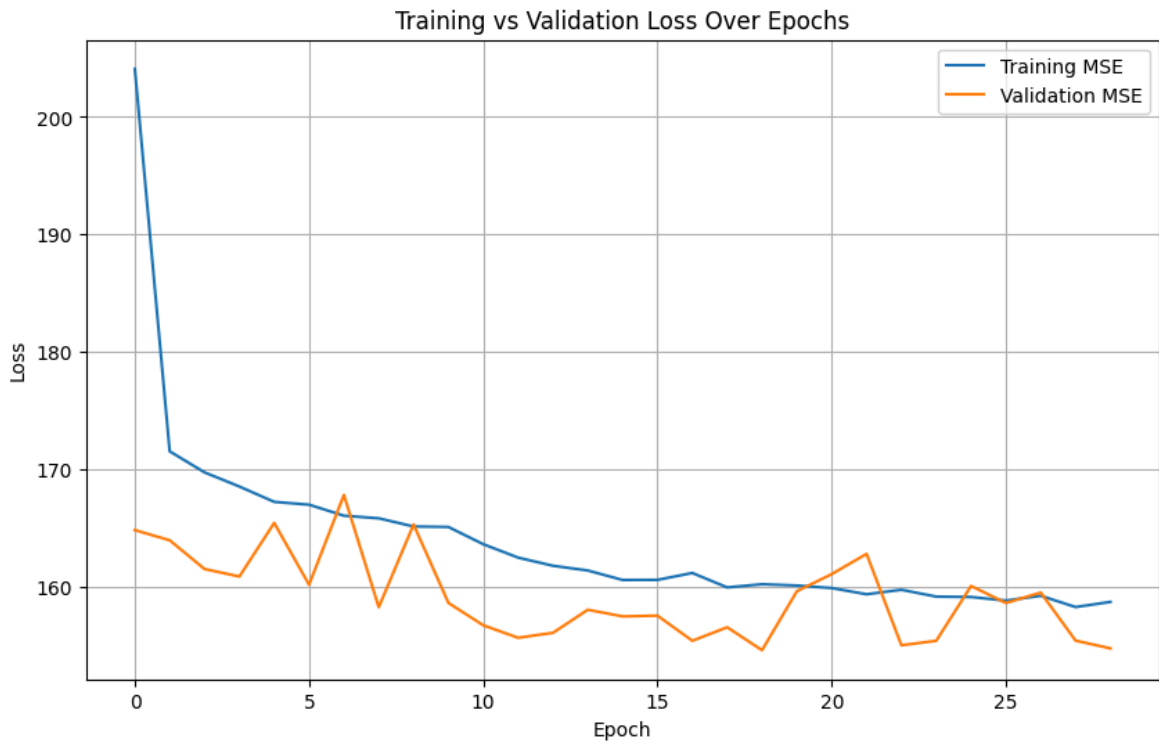
Epoch 20/50
1421/1421  **13s** 9ms/step - loss: 159.6297 - mae: 10.0076 - val_loss: 159.6281 - val_mae: 9.8194

Epoch 21/50
1421/1421 ————— 13s 9ms/step - loss: 160.2094 - mae: 10.0259 - val_loss: 161.0898 - val_mae: 9.8277
Epoch 22/50
1421/1421 ————— 12s 9ms/step - loss: 159.7376 - mae: 10.0306 - val_loss: 162.8111 - val_mae: 9.8998
Epoch 23/50
1421/1421 ————— 14s 10ms/step - loss: 158.6777 - mae: 9.9571 - val_loss: 155.0361 - val_mae: 9.9213
Epoch 24/50
1421/1421 ————— 15s 11ms/step - loss: 159.2619 - mae: 9.9726 - val_loss: 155.4119 - val_mae: 9.9142
Epoch 25/50
1421/1421 ————— 13s 9ms/step - loss: 158.9671 - mae: 9.9737 - val_loss: 160.0805 - val_mae: 9.7909
Epoch 26/50
1421/1421 ————— 14s 10ms/step - loss: 159.3607 - mae: 9.9716 - val_loss: 158.6290 - val_mae: 9.7932
Epoch 27/50
1421/1421 ————— 14s 10ms/step - loss: 158.5242 - mae: 9.9351 - val_loss: 159.5169 - val_mae: 9.8321
Epoch 28/50
1421/1421 ————— 14s 10ms/step - loss: 157.4952 - mae: 9.9380 - val_loss: 155.4253 - val_mae: 9.8607
Epoch 29/50
1421/1421 ————— 13s 9ms/step - loss: 160.1654 - mae: 10.0049 - val_loss: 154.7669 - val_mae: 9.7804

```
In [49]: fig, ax = plt.subplots(figsize=(10, 6))
ax.plot(history.history['loss'], label='Training MSE')
ax.plot(history.history['val_loss'], label='Validation MSE')

ax.set_title('Training vs Validation Loss Over Epochs')
ax.set_xlabel('Epoch')
ax.set_ylabel('Loss')

ax.legend(loc='best')
plt.grid(True)
plt.show()
```



```
In [50]: y_pred = model.predict(X_test)
print("MSE:", mean_squared_error(y_test, y_pred)) #MSE
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred))) #RMSE
print("MAE :", mean_absolute_error(y_test, y_pred) ) #MAE
print("MAPE :", mean_absolute_percentage_error(y_test, y_pred) ) #MAPE
print("r2_score:", r2_score(y_test, y_pred)) #r2score
```

888/888 ————— 4s 4ms/step

MSE: 155.6469484899873

RMSE: 12.475854619623751

MAE : 9.926205313844445

MAPE : 0.25493256842432643

r2_score: 0.22020153682583432

Conclusion:

- The Logistic Regression model required considerably more time for training compared to the Neural Network model and Random Forest Regressor.
- Prediction time was also longer with the Neural Network model.
- Generally, neural networks are known to deliver higher accuracy than random forest regressors across various problem types, as they are capable of learning more intricate relationships between input features and target variables.

Questions

1. Defining the problem statements and where can this and modifications of this be used?

- Porter, India's largest marketplace for intra-city logistics, partners with restaurants to handle food deliveries.

- The company aims to provide its customers with accurate delivery time estimates based on several influencing factors.
 - Key factors include order time, items ordered, delivery partner availability, the number of outstanding orders, and the type of restaurant.
 - The objective is to develop a regression model that can predict delivery times with high accuracy using these features.
-

2. List 3 functions the pandas datetime provides with one line explanation.

- **pd.to_datetime():** Converts input data to datetime format.
 - **pd.date_range():** Creates a sequence of dates between specified start and end.
 - **pd.to_timedelta:** Converts input to a timedelta, representing durations.
-

3. Short note on datetime, timedelta, time span (period)

- **datetime:** Represents a specific point in time, including date and time components. It is used for precise timekeeping and manipulating time-related data.
 - **timedelta:** Represents a duration or difference between two dates or times. It is used for arithmetic operations involving dates and times, such as adding or subtracting durations.
 - **time Span (Period):** Defines a range or interval of time. It can be represented by a pair of datetime objects (start and end) or by a datetime object combined with a timedelta to specify the extent of the period.
-

4. Why do we need to check for outliers in our data?

- Outliers are data points that deviate significantly from the majority of the dataset. They can arise from various sources, such as data entry errors, and may distort statistical analyses, making it challenging to draw precise conclusions.
 - Here's why identifying outliers is crucial:
 1. **Impact on Statistical Analysis:** Outliers can disproportionately affect statistical measures like the mean, median, and standard deviation. They may skew these values, leading to misleading results and inaccurate interpretations.
 2. **Challenges in Drawing Conclusions:** Outliers can obscure the true distribution of data and mask underlying patterns. This distortion makes it harder to derive meaningful insights and understand the data accurately.
 3. **Errors in Data Entry:** Outliers might be introduced due to mistakes in data recording. For instance, an incorrectly entered value can create an outlier, so it's essential to detect and address these anomalies before performing data analysis.
-

5. Name 3 outlier removal methods?

- Here are three common methods for removing outliers:
 1. **Z-Score Method:** Identifies outliers by calculating the Z-score, which measures how many standard deviations a data point is from the mean. Data points with a Z-score beyond a certain threshold are considered outliers.
 2. **IQR Method:** Uses the Interquartile Range (IQR), which is the range between the first (25th percentile) and third quartile (75th percentile). Data points that fall below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$ are flagged as outliers.
 3. **Isolation Forest:** A machine learning-based method that isolates observations by randomly selecting features and splitting values. Outliers are identified as observations that are isolated earlier in the process compared to normal data points.
-

6. What classical machine learning methods can we use for this problem?

- For outlier detection, classical machine learning methods include:
 1. **K-Nearest Neighbors (KNN):** Uses distance metrics to find the k-nearest neighbors of a data point. Outliers are identified as points that have a larger average distance to their neighbors compared to the majority of points.
 2. **One-Class SVM (Support Vector Machine):** Trains a model to separate the data from the origin in a high-dimensional space. Data points that fall outside the learned boundary are considered outliers.
 3. **Isolation Forest:** Specifically designed for anomaly detection, it isolates observations by randomly selecting features and splitting values. Points that are isolated quickly are likely outliers.
 4. **Local Outlier Factor (LOF):** Measures the local density deviation of a data point compared to its neighbors. Points with significantly lower density compared to their neighbors are flagged as outliers.
 5. **Robust PCA (Principal Component Analysis):** Decomposes the data into a low-rank matrix and a sparse matrix. Outliers are captured in the sparse matrix, making it possible to detect anomalous data points.
-

7. Why is scaling required for neural networks?

- Scaling is crucial for neural networks due to several reasons:
 1. **Improves Convergence:** Scaling features to a similar range (e.g., [0, 1] or standardizing to zero mean and unit variance) helps the neural network converge faster during training. It ensures that gradients are more stable and prevents some features from disproportionately influencing the weight updates.

2. **Reduces Numerical Instability:** Neural networks often involve operations that can lead to numerical instability, such as large matrix multiplications and activation functions. Scaling inputs helps mitigate issues related to overflow or underflow, ensuring more stable computations.
 3. **Facilitates Activation Function Efficiency:** Many activation functions (like sigmoid and tanh) are sensitive to the scale of the inputs. Scaling ensures that inputs to these functions lie within the region where they can effectively learn and provide meaningful gradients.
 4. **Balances Feature Importance:** Different features might have different scales (e.g., height in cm vs. weight in kg). Scaling ensures that each feature contributes equally to the model's learning process, preventing features with larger scales from dominating the learning process.
 5. **Improves Training Speed:** When features are on a similar scale, the optimization process (e.g., gradient descent) can proceed more smoothly and efficiently, potentially reducing the number of epochs required to train the model.
-

8. Briefly explain your choice of optimizer.

- **Combines Momentum and RMSProp:** Adam integrates the benefits of both Momentum (which uses moving averages of past gradients) and RMSProp (which uses moving averages of squared gradients).
 - **Adaptive Learning Rates:** It automatically adjusts learning rates for each parameter, allowing for efficient and reliable convergence.
 - **Faster Convergence:** The combination of momentum and adaptive learning rates often leads to faster and more stable convergence compared to traditional optimizers.
 - **Robust Performance:** Adam performs well across a wide range of problems and is less sensitive to hyperparameter tuning.
 - **Efficient:** It is computationally efficient and has low memory requirements, making it suitable for large-scale machine learning problems.
-

9. Which activation function did you use and why?

1. ReLU: A Popular Activation Function

- ReLU, widely used in neural networks and deep learning, addresses the vanishing gradient problem seen in functions like sigmoid and tanh. By introducing sparsity, ReLU deactivates neurons with negative inputs, improving network efficiency.

2. Advantages of ReLU in Neural Networks

- ReLU is a favored activation function in deep learning for its ability to tackle the vanishing gradient problem. It introduces sparsity by deactivating neurons with

negative inputs, enhancing the performance of neural networks.

3. ReLU's Benefits in Deep Learning

- ReLU is commonly used in neural networks because it helps reduce the vanishing gradient problem, unlike sigmoid and tanh. It also creates sparse activation by turning off neurons with negative outputs, making the model more efficient.

10. Why does a neural network perform well on a large dataset?

- A neural network performs well on a large dataset because larger datasets provide more diverse and comprehensive information for the model to learn from, leading to better generalization. Here are the key reasons:
 1. Better Learning of Patterns: Large datasets capture a wider range of patterns, relationships, and nuances in the data, allowing the neural network to learn complex features effectively.
 2. Reduced Overfitting: With more data, the model is less likely to memorize specific samples, reducing the risk of overfitting. This leads to a model that performs better on unseen data.
 3. Improved Generalization: Large datasets provide the network with a broader understanding of the underlying distribution of the data, enhancing its ability to make accurate predictions on new inputs.
 4. Balanced Classes: Larger datasets are more likely to provide balanced class distributions, which helps the model learn equally well across different classes, improving overall performance.
 5. Enhanced Training Stability: With more examples, the gradient estimates during training are more accurate, leading to smoother and more stable convergence.
- Overall, a large dataset equips neural networks with the information needed to learn complex decision boundaries and improve their predictive capabilities.

In []: